

**CNN**

Step 1: Assuming we are having a  
6x6 data:

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

1: wet  
2: dry

Step 1: Assuming we are having a 6x6 data:

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

1: wet  
2: dry

Step 2: the neuron (or filter) in CNN includes some basic information of the data ~ a filter is a matrix, e.g., the following example gives two 3x3 filters

Filter 1			Filter 2		
1	-1	-1	-1	1	-1
-1	1	-1	-1	1	-1
-1	-1	1	-1	1	-1

- The elements in each filter is the parameters to be determined through the training process (like the weights in ANN).
- in this example, if we have a 3x3 filter, then CNN only look at the area of 3x3 in the image at once instead of every grid point.

Step 1: Assuming we are having a 6x6 data:

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

1: wet  
2: dry

Step 2: the neuron (or filter) in CNN includes some basic information of the data ~ a filter is a matrix, e.g., the following example gives two 3x3 filters

Filter 1			Filter 2		
1	-1	-1	-1	1	-1
-1	1	-1	-1	1	-1
-1	-1	1	-1	1	-1

- The elements in each filter is the parameters to be determined through the training process (like the weights in ANN).
- in this example, if we have a 3x3 filter, then CNN only look at the area of 3x3 in the image at once instead of every grid point.

Step 3: Applying the Filter onto the image from the top-left corner:

1	-1	-1	1	0	0	0	0	0
-1	1	-1	0	1	0	0	0	0
-1	-1	1	0	0	1	1	0	0
			0	0	0	1	1	0
			0	0	0	0	1	0
			0	0	0	0	1	0

Step 3: Applying the Filter  
onto the image from the top-  
left corner:

1	-1	-1	1	0	0	0	0	0	0
-1	1	-1	0	1	0	0	0	0	0
-1	-1	1	0	0	1	1	0	0	0
			0	0	0	1	1	0	0
			0	0	0	0	1	0	0
			0	0	0	0	1	0	0

Step 4: Generating the inner  
product between the filter  
the corresponding area of  
the image

1	-1	-1
-1	1	-1
-1	-1	1

 • 

1	0	0
0	1	0
0	0	1

 = 3

Inner product

The diagram shows a 3x3 submatrix of the distance matrix, highlighted with a red box. This submatrix represents the distances between nodes 1, 2, and 3, considering only nodes 1, 2, and 3 as potential intermediate nodes. The values are:

1	-1	-1
-1	1	-1
-1	-1	1

Red lines connect this submatrix to the corresponding rows and columns in the larger 6x6 distance matrix, showing how the algorithm updates the shortest paths by considering the third node as an intermediate node.

$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = 3$$

Inner product

Diagram illustrating Stride=1 and Stride=2 for a 2D convolution. The input grid is 7x7, and the kernel is 3x3. The output grid is 5x5. The kernel is highlighted in red for Stride=1 and green for Stride=2. The output grid shows the kernel's output values highlighted in blue.

	Stride=1			Stride=2		
1	0	0	0	0	0	0
0	1	0	0	0	0	0
0	0	1	1	0	0	0
0	0	0	1	1	0	0
0	0	0	0	1	0	0
0	0	0	0	0	1	0
0	0	0	0	0	1	0

Step 3: Applying the Filter onto the image from the top-left corner:

1	-1	-1	1	0	0	0	0	0	0
-1	1	-1	0	1	0	0	0	0	0
-1	-1	1	0	0	1	1	0	0	0
			0	0	0	1	1	0	0
			0	0	0	0	1	0	0
			0	0	0	0	1	0	0

Step 3.1: Generating the inner product between the filter the corresponding area of the image

1	-1	-1	1	0	0	= 3
-1	1	-1	0	1	0	
-1	-1	1	0	0	1	

Inner product

Step 4: moving the filter the entire dataset based on a predefined “stride”, and get the inner product for each movement

		Stride=1	Stride=2								
1	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0
0	0	1	1	0	0	0	0	0	0	0	0
0	0	0	1	1	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0

the purpose of this filter is to detect if there is similar pattern of the filter in the image. The higher number in the inner product matrix usually means the higher similarity between the filter and the corresponding area of image

Note that in a CNN there are usually a dozen of filters.

Obtained from Step 2	Obtained from Step 4										
3	-1	-3	-1	1	0	0	0	0	0	0	0
-3	1	0	-3	0	1	0	0	0	0	0	0
-3	-3	0	1	0	0	1	1	0	0	0	0
3	-2	-2	-1	0	0	0	1	0	0	0	0

The last filtering process

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

↑  
Input data matrix

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

4x4 data

↑  
The inner product  
matrix from a filter



1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

↑  
Input data matrix

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

4x4 data

↑  
The inner product  
matrix from a filter

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

pooling

↑  
Using "pooling" to  
reduce the size of  
the inner product

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

↑  
Input data matrix

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

4x4 data

↑  
The inner product  
matrix from a filter

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

pooling

↑  
Using "pooling" to  
reduce the size of  
the inner product

Taking the maximum  
value from this  
subdomain

3	0
3	0
3	0

2x2 data

We can have many of  
these small matrices  
(depth = the number  
of filters)

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

1	1
2	0
3	0
4	0
5	0

7	0
8	1
9	0
	0
	0

13	0
14	0
15	1
	0

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1



1	0	0	0	0	0
0	1	0	0	0	0
0	0	1	1	0	0
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

1	1
2	0
3	0
4	0
5	0

7	0
8	1
9	0
	0
	0

13	0
14	0
15	1
	0

3	1	-3	-1
2	1	0	-3
3	-3	0	1
3	-2	-2	-1

location of the item

1	2	3			
7	8	9			
1 3	1 4	1 5			
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data

1	1
2	0
3	0
4	0
5	0

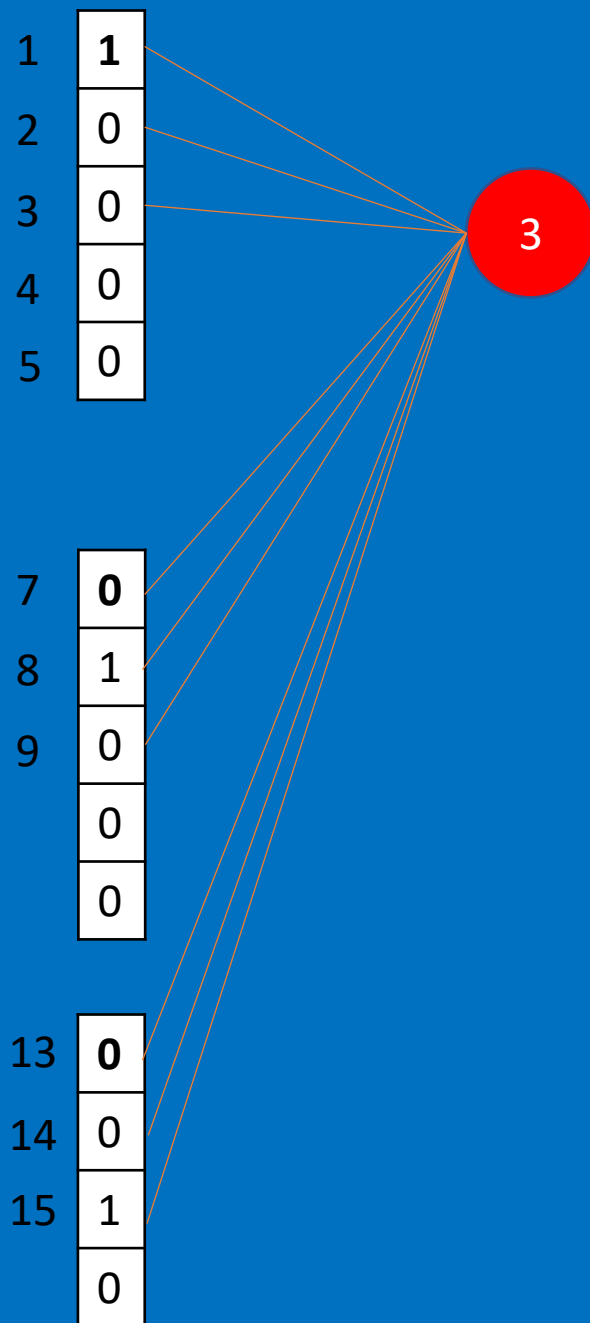
7	0
8	1
9	0
	0
	0

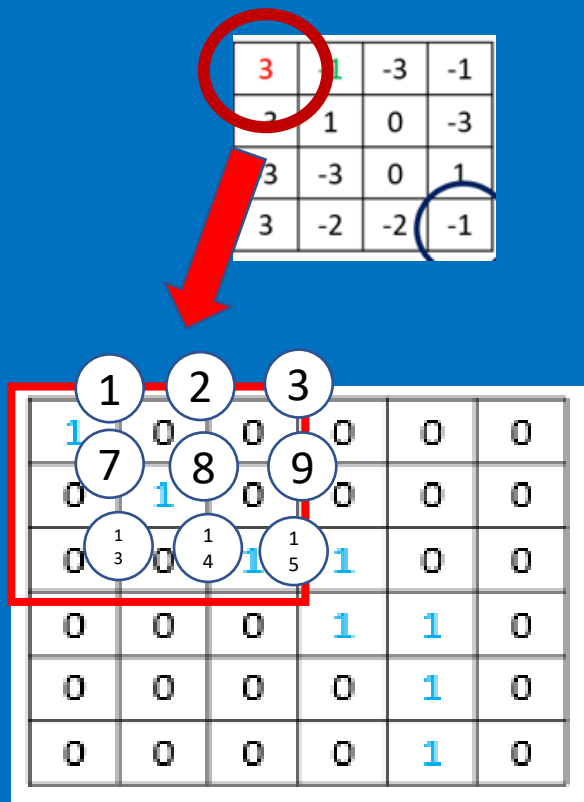
13	0
14	0
15	1
	0

3	1	-3	-1
2	1	0	-3
3	-3	0	1
3	-2	-2	-1

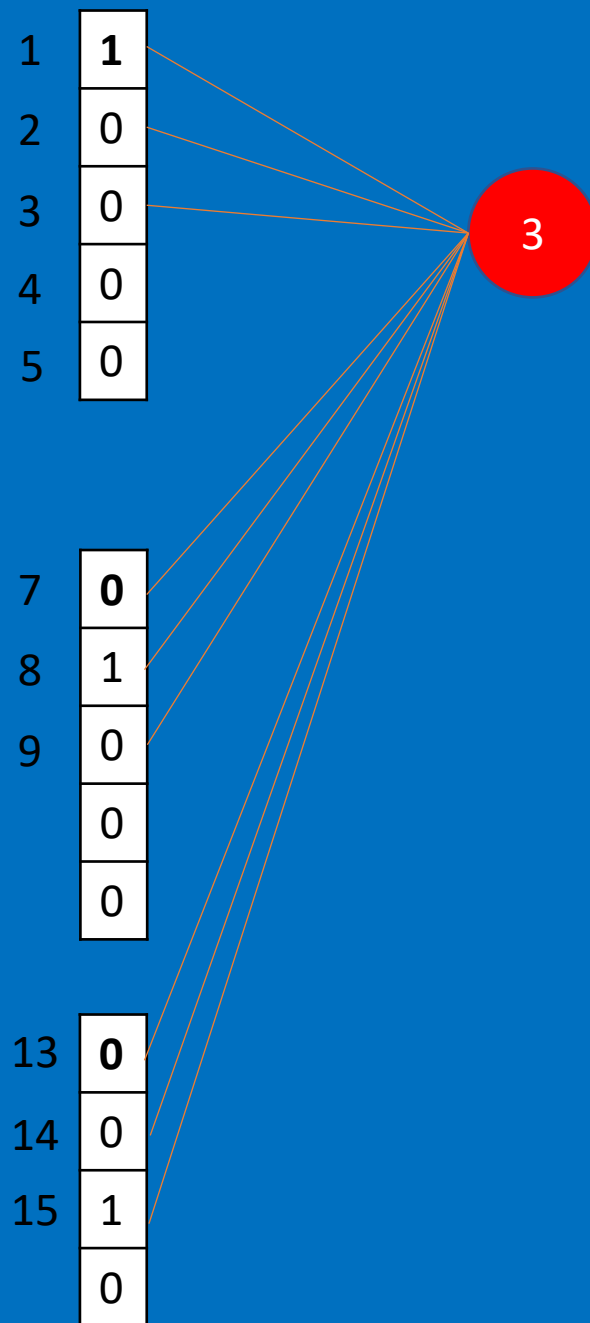
1	2	3			
1	7	8	9		
0	1	0	0		
0	1	4	5	1	
0	0	0	1	1	0
0	0	0	0	1	0
0	0	0	0	1	0

6x6 data



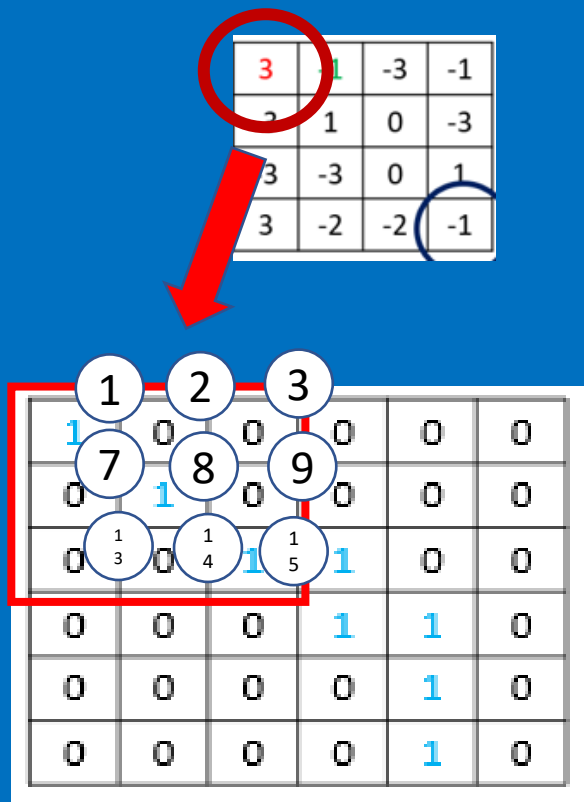


6x6 data

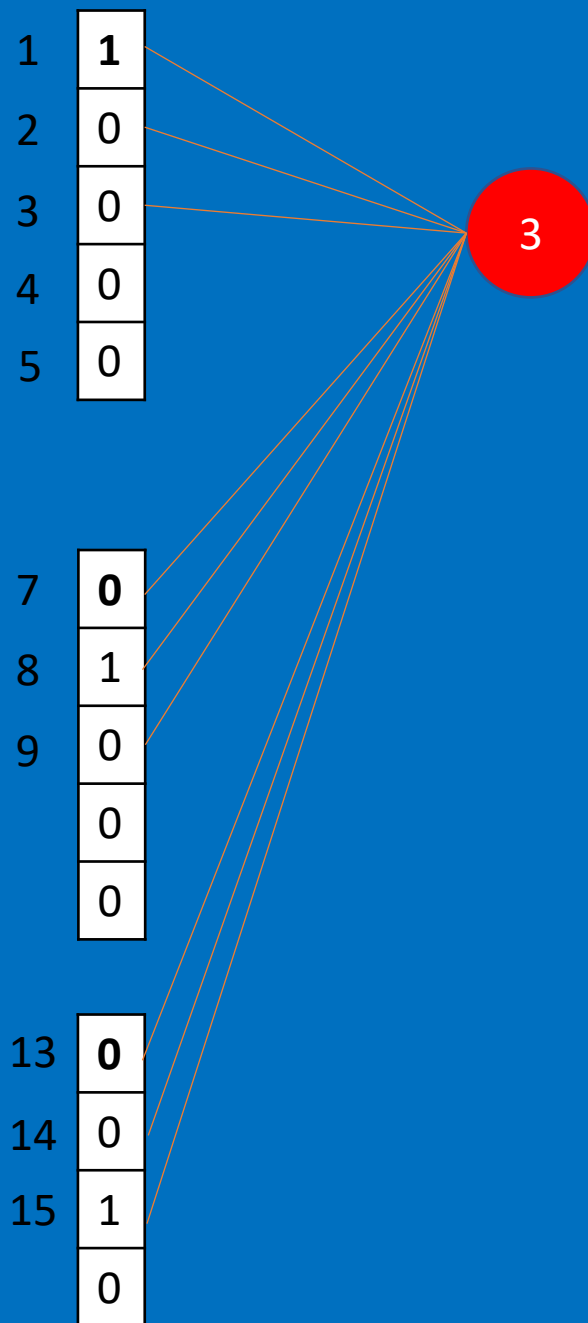


So for this neuron ("3") it only needs to have the weights for some of the elements





6x6 data



So for this neuron ("3") it only needs to have the weights for some of the elements

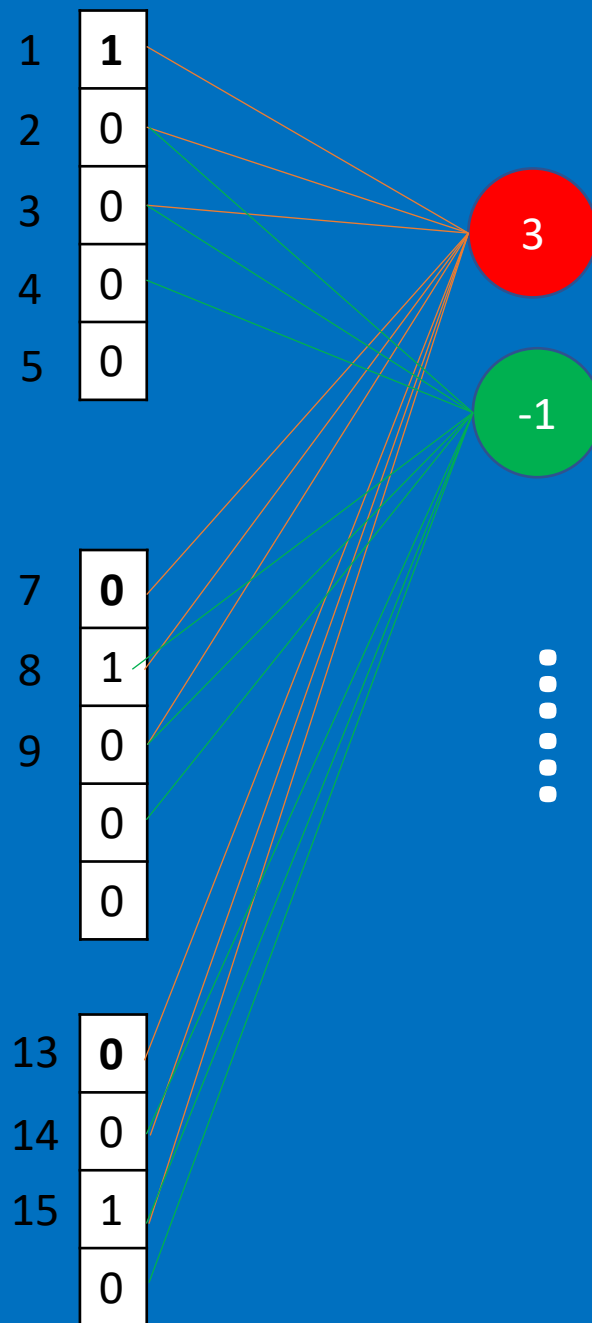
1	-1	-1
-1	1	-1
-1	-1	1

And the weight for these 9 synapses is "defined" by the filter itself

3	-1	-1	-1
-3	1	0	-3
-3		0	1
3	2	-2	-1

	2	3	4	
1	0	8	9	1
0	1	0	0	0
0	0	1	5	6
0	0	0	1	1
0	0	0	0	1

6x6 data

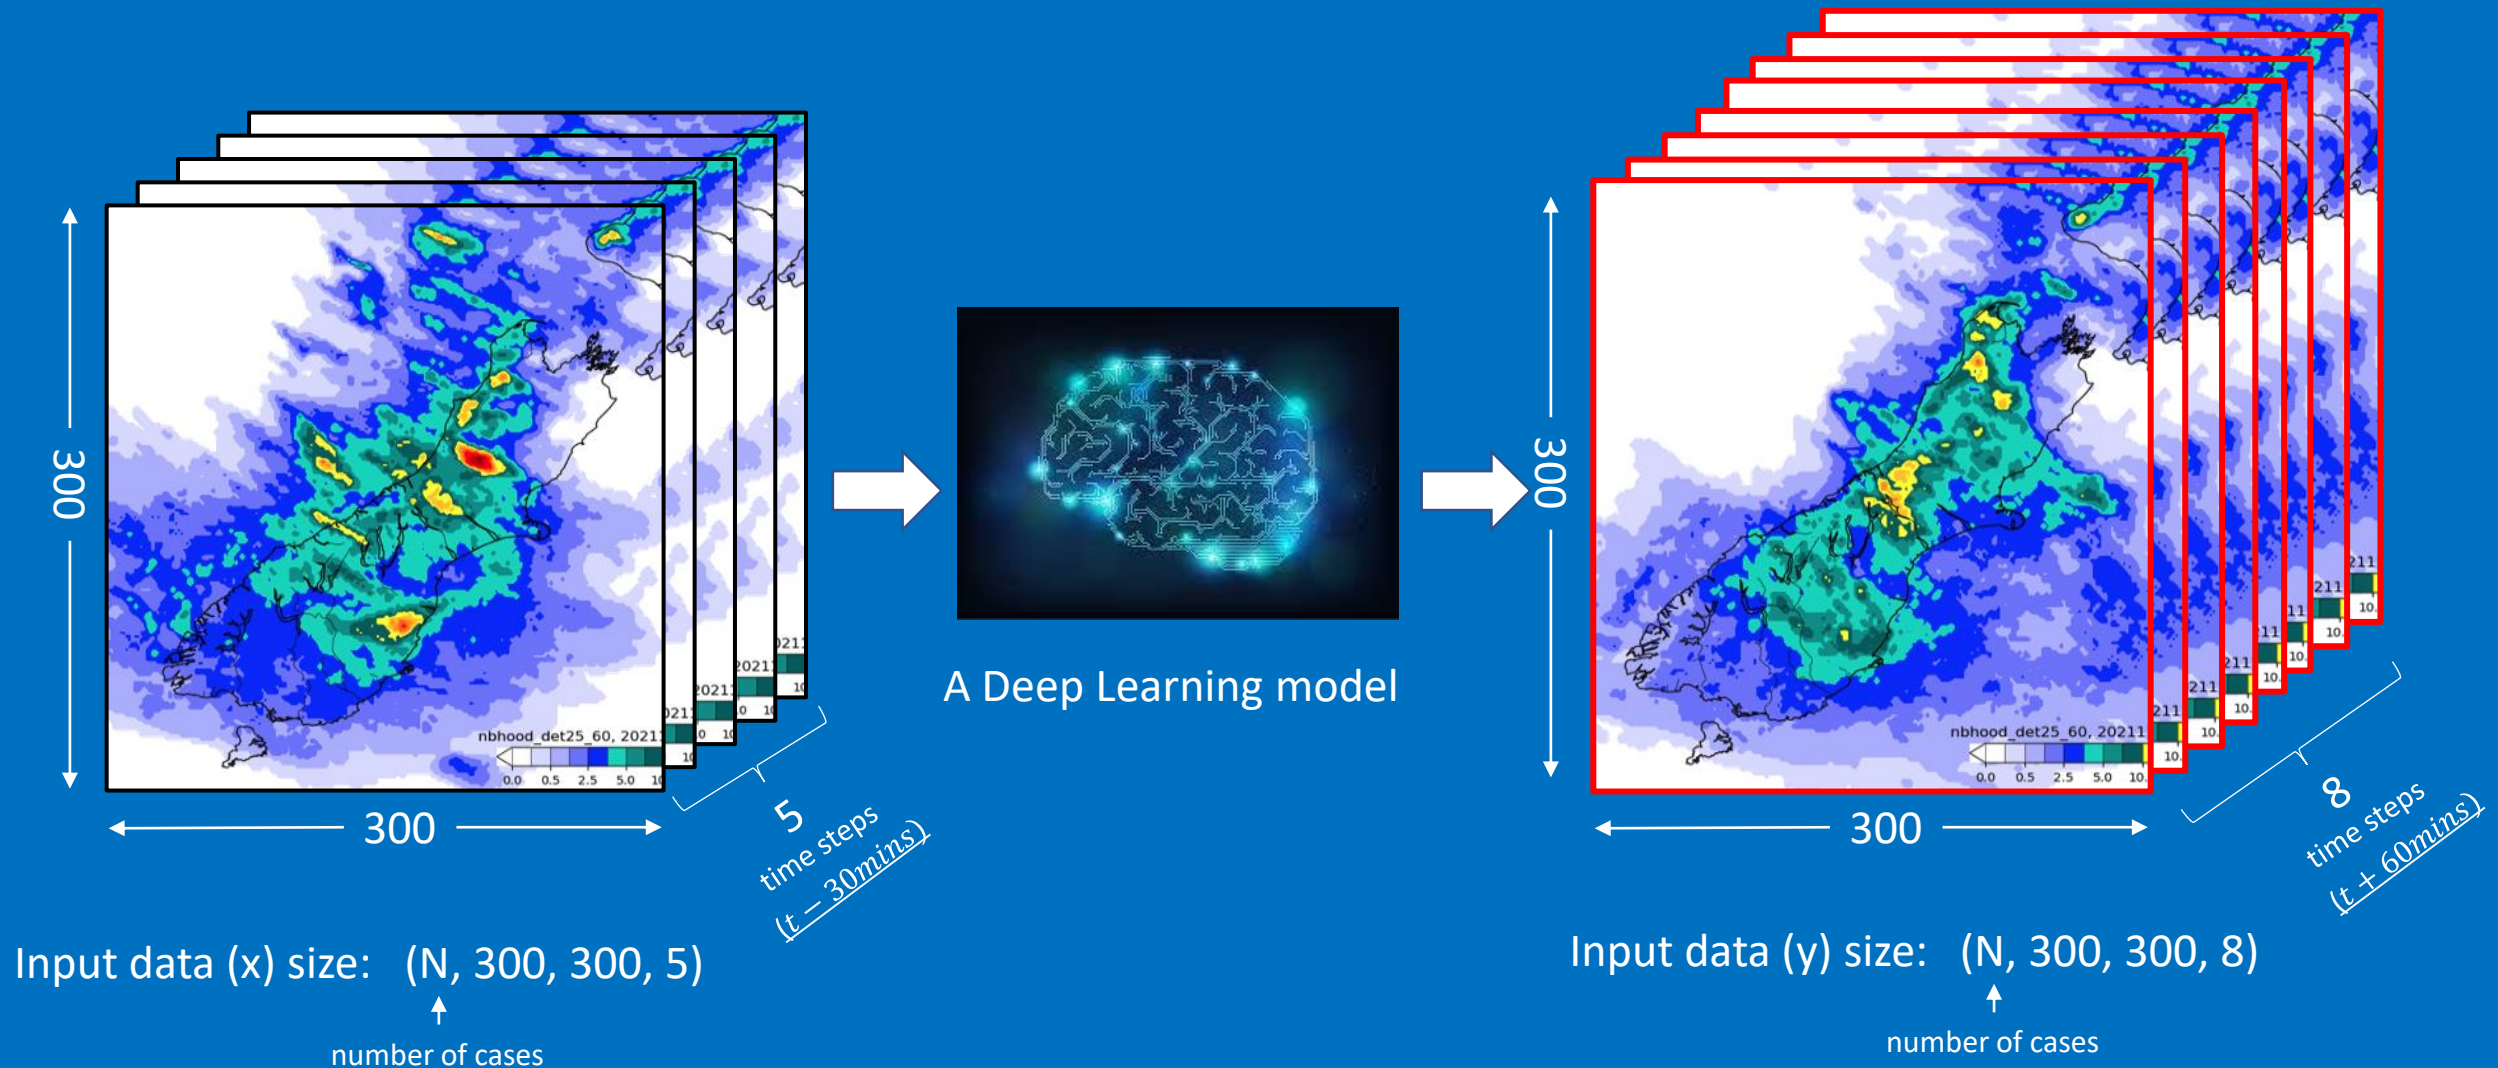


Therefore, the solution of CNN turns to something like ANN, all we need to do is just to solve the weights (which included in the filter) through the training (e.g., the cost function in ANN).

# Best practice for CNN

<https://towardsdatascience.com/a-guide-to-an-efficient-way-to-build-neural-network-architectures-part-ii-hyper-parameter-42efca01e5d7>

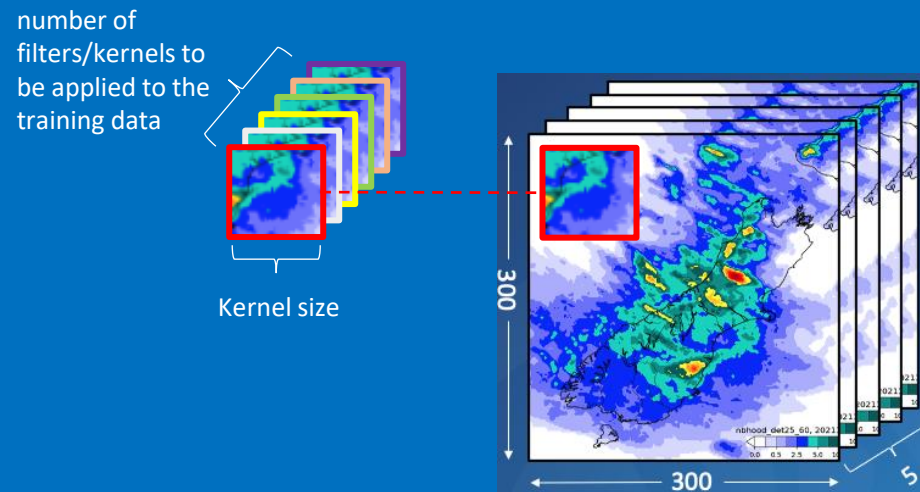
<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>



# Best practice for CNN: How filter works?

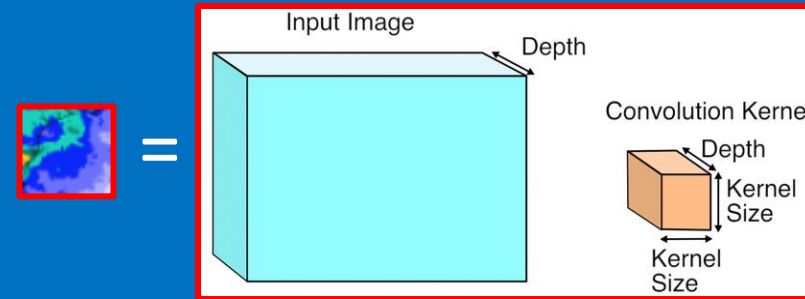
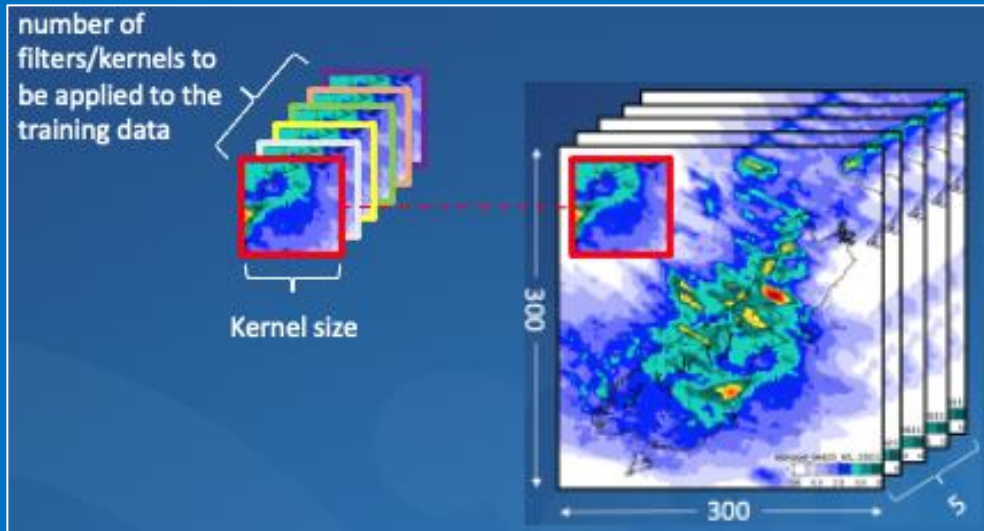
<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>

<https://stackoverflow.com/questions/43306323/keras-conv2d-and-input-channels>



# Best practice for CNN: How filter works?

<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>  
<https://stackoverflow.com/questions/43306323/keras-conv2d-and-input-channels>



in Conv2D, a filter actually is 3D which contains the depth of the image (in this case the depth is 5)

<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>  
<https://stackoverflow.com/questions/43306323/keras-conv2d-and-input-channels>

The diagram shows a large light blue rectangular prism representing the 'Input Image'. It has three dimensions: width, height, and depth. An arrow labeled 'Depth' points to the right side of the prism. To the right of the input image is a smaller orange rectangular prism representing the 'Convolution Kernel'. It also has three dimensions: width, height, and depth. Arrows labeled 'Depth', 'Kernel Size', and 'Kernel Size' point to the depth, height, and width of the kernel respectively.

## Inner product combination

Put the combined inner product into the “inner product matrix”

$$X = X_1 + X_2 + X_3 + X_4 + X_5$$

A 6x6 grid is shown. The top-left cell contains a red 'X'. A red arrow points from the 'X' to the top-left corner of the grid. The grid is part of a larger blue area on the left.

# Best practice for CNN: How filter works?

<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>  
<https://stackoverflow.com/questions/43306323/keras-conv2d-and-input-channels>

## Thumb of rules: How to set the number of filters

In most cases, over layers the number of filters is generally ascending. The reason is that:

The higher the number of filters, the higher the number of *abstractions* that your Network can extract from image data.

The reason why the number of filters is generally ascending is that at the input layer the Network receives raw pixel data. Raw data are always noisy, and this is especially true for image data.

Because of this, we let CNNs extract first some relevant information from noisy, "dirty" raw pixel data. Once the useful features have been extracted, then we make the CNN elaborate more complex abstractions on it.

That is why the number of filters usually increases as the Network gets deeper, even though it doesn't necessarily have to be like that.



# Best practice for CNN: Why max pooling

## Reducing Computational Load

Since max pooling is reducing the resolution of the given output of a convolutional layer, the network will be looking at larger areas of the image at a time going forward, which reduces the amount of parameters in the network and consequently reduces computational load.

## Reducing Overfitting

Additionally, max pooling may also help to reduce overfitting.

For example, from a radar image, a filter looks for the pattern which is closest to itself. From the output of the convolutional layer, we can think of the higher valued pixels as being the ones that are the most similar to the filter.

With max pooling, as we're going over each region from the convolutional output, we're able to pick out the most activated pixels and therefore “only” preserve the features match the filter the most while discarding those features which are not very close to the filter

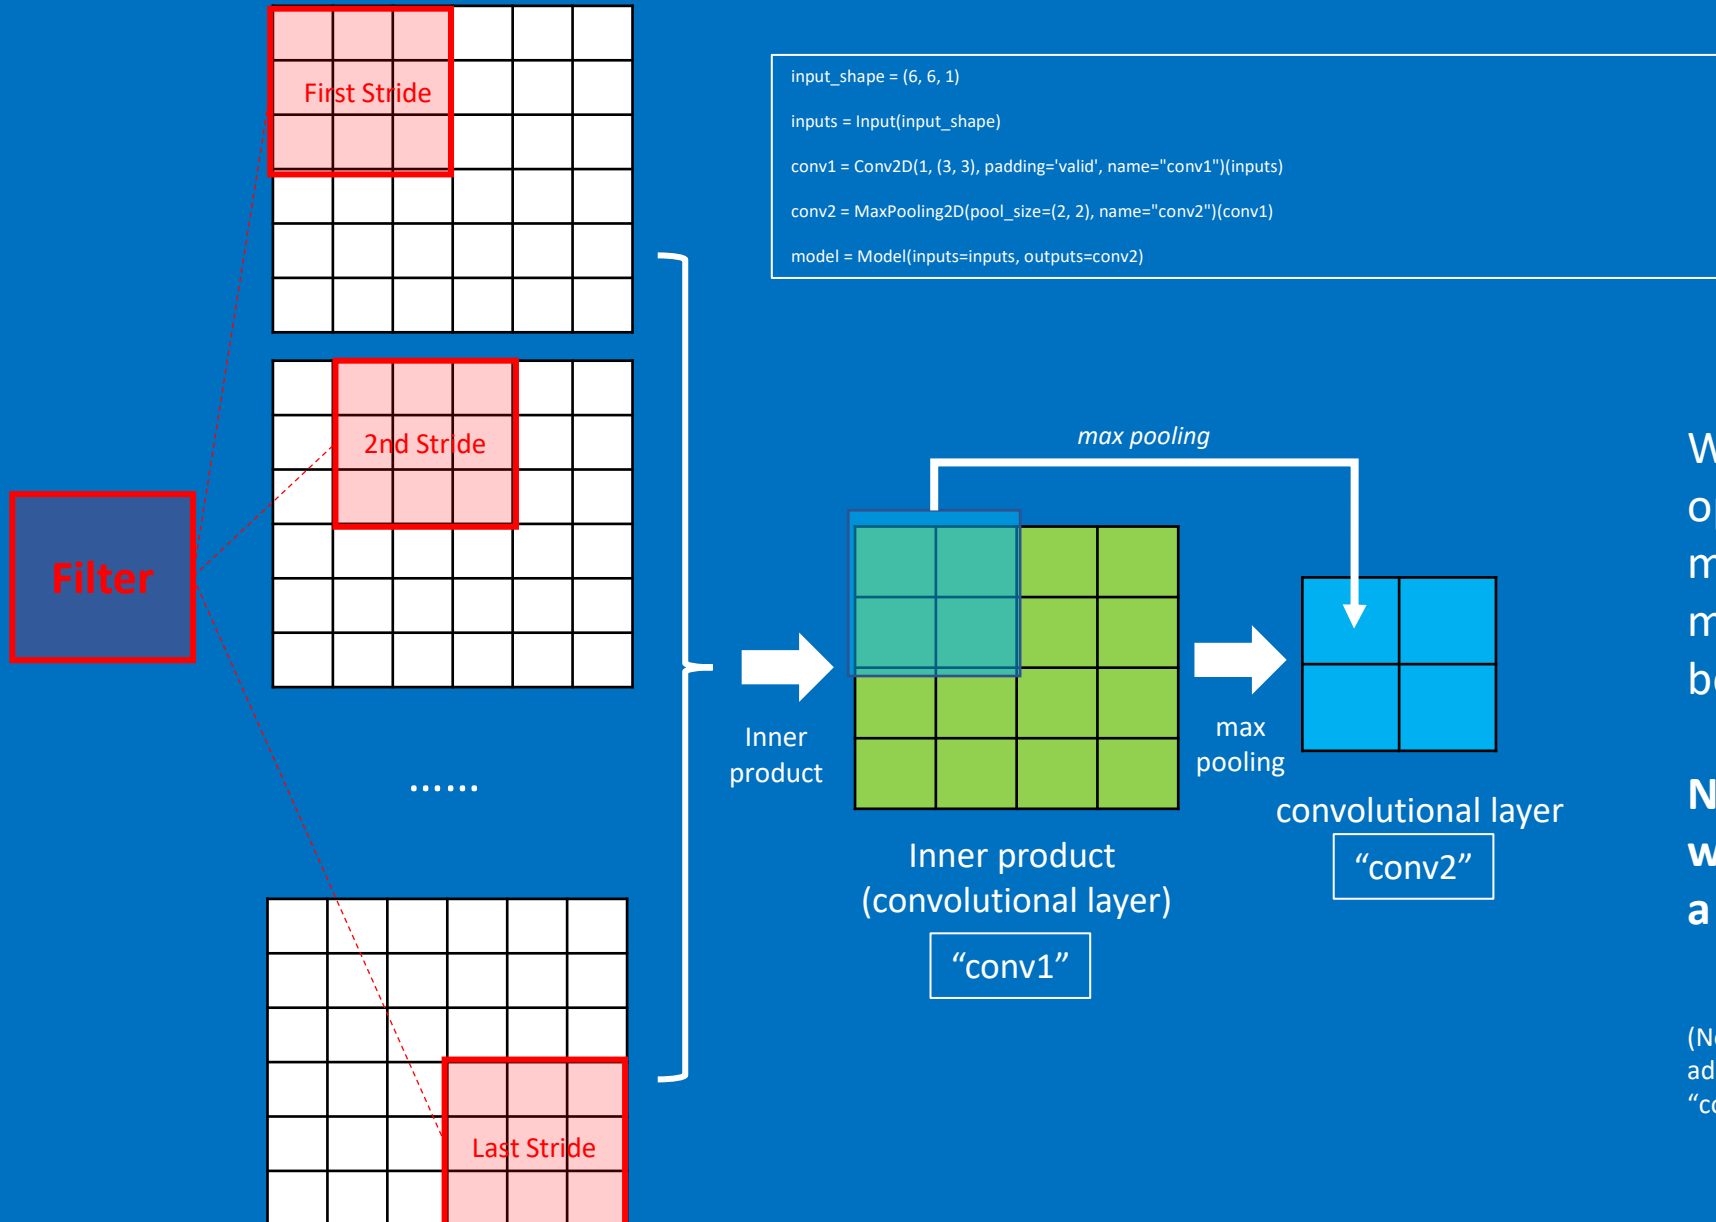
## Drop out

We randomly remove neurons in the neural network during the training to reduce the chances of overfitting



# Best practice for CNN: up-sampling & conv transpose

<https://towardsdatascience.com/understand-transposed-convolutions-and-build-your-own-transposed-convolution-layer-from-scratch-4f5d97b2967>



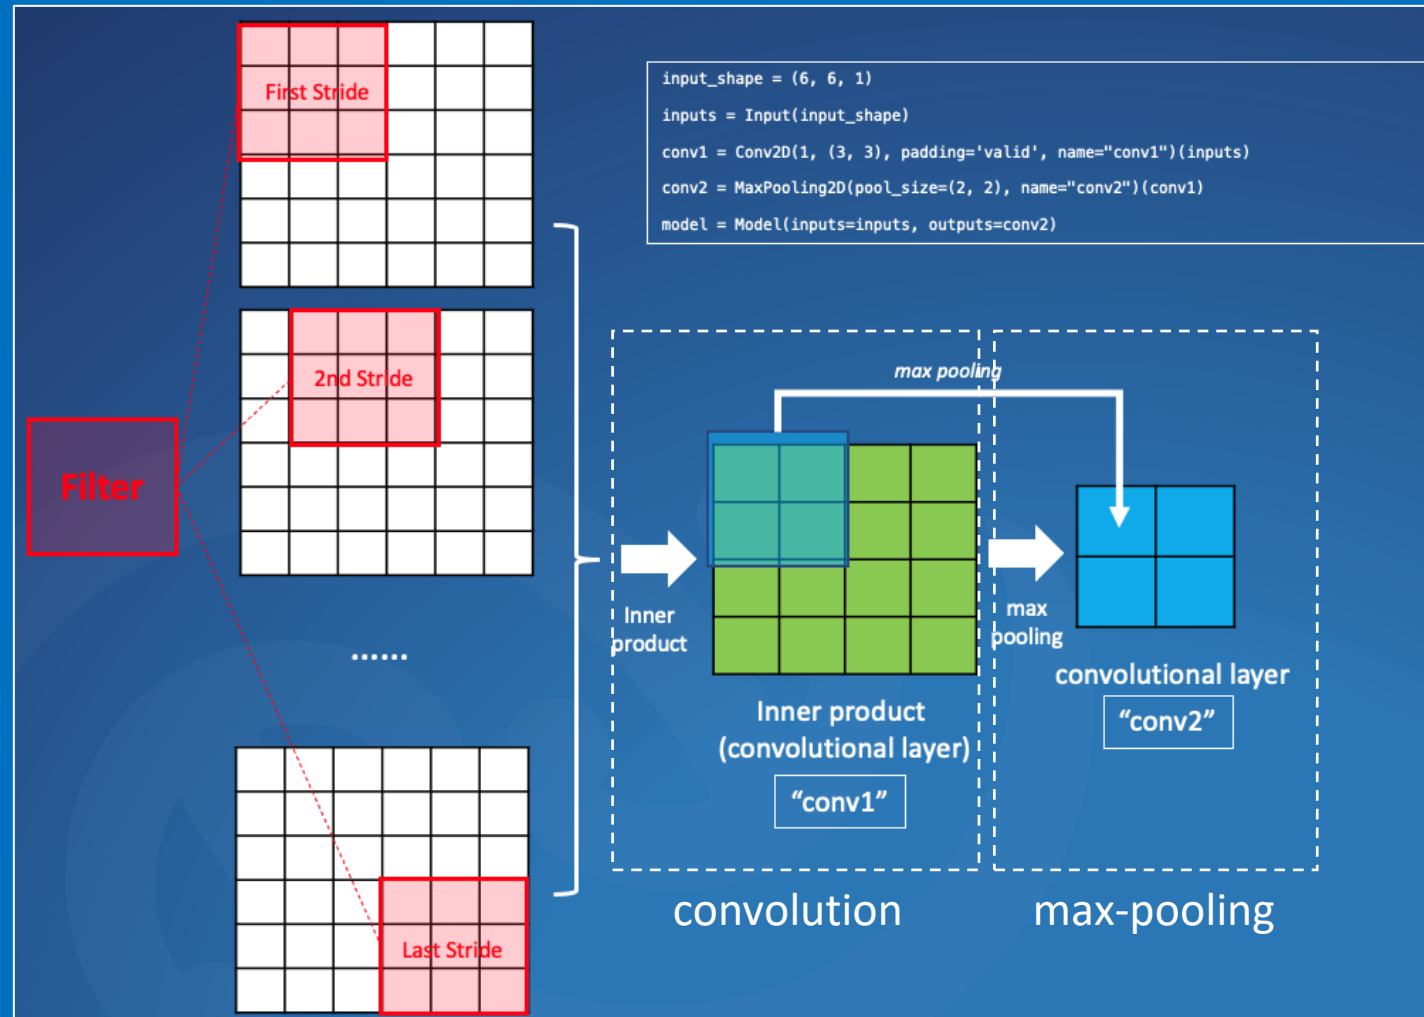
Without any padding, this operation transforms a “6x6” matrix into a “4x4” matrix, after max-pooling the matrix becomes “2x2”

**Now, our question is: what if we want to go backward from a 2x2 matrix to a 6x6 matrix?**

(Note that an additional “padding” can be used to address the shape issue for “conv1”, but not for “conv2”)

# Best practice for CNN: up-sampling & conv transpose

<https://towardsdatascience.com/understand-transposed-convolutions-and-build-your-own-transposed-convolution-layer-from-scratch-4f5d97b2967>



In order to restore the shape from “2x2” to “6x6”, first we need to understand that two activities made the shape changes:

- convolution
- max-pooling

In order to resume the shape from

- convolution: we can use padding or conv-transposed
- max-pooling: we can use up-sampling

# Best practice for CNN: up-sampling & conv transpose

<https://towardsdatascience.com/understand-transposed-convolutions-and-build-your-own-transposed-convolution-layer-from-scratch-4f5d97b2967>

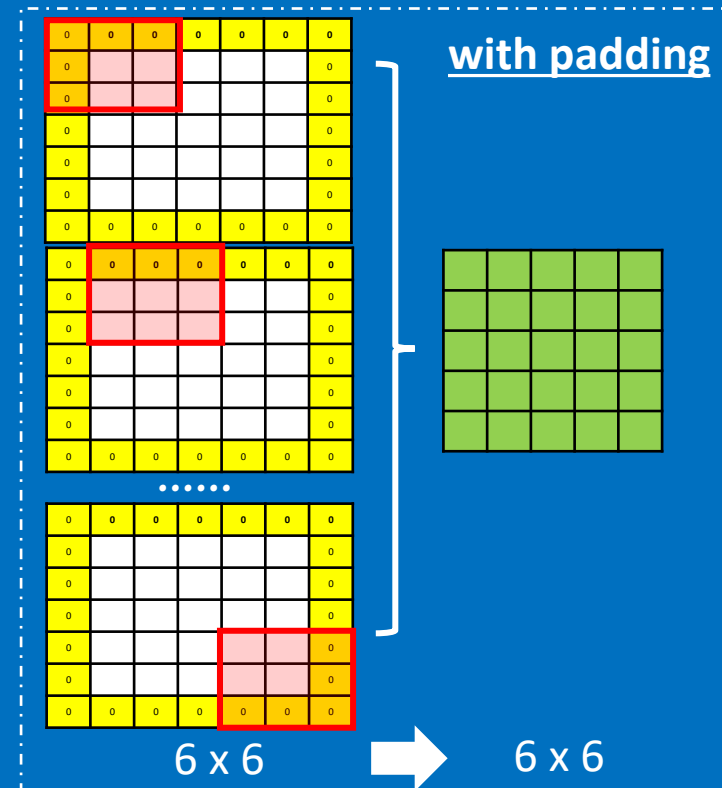
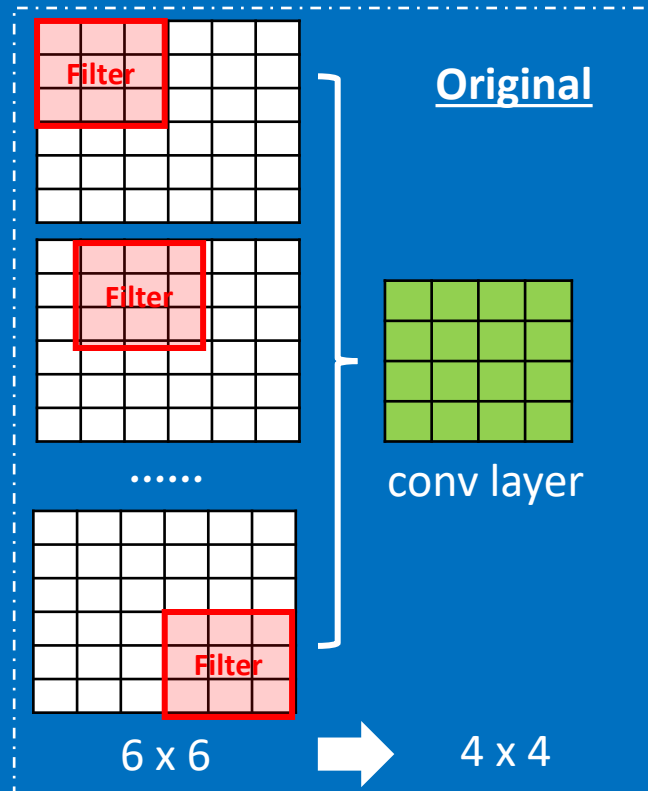
## What is padding

In order to resume the shape from

- **convolution:** we can use **padding**, or conv-transposed
- **max-pooling:** we can use up-sampling or conv-transpose

As we just discussed, the convolutional layers reduce the size of the output.

Padding basically extends the area of an image. The kernel/filter which moves across the image scans each pixel and converts the image into a smaller image. In order to work the kernel with processing in the image, padding is added to the outer frame of the image to allow for more space for the filter to cover in the image. Adding padding to an image processed by a CNN allows for a more accurate analysis of images.



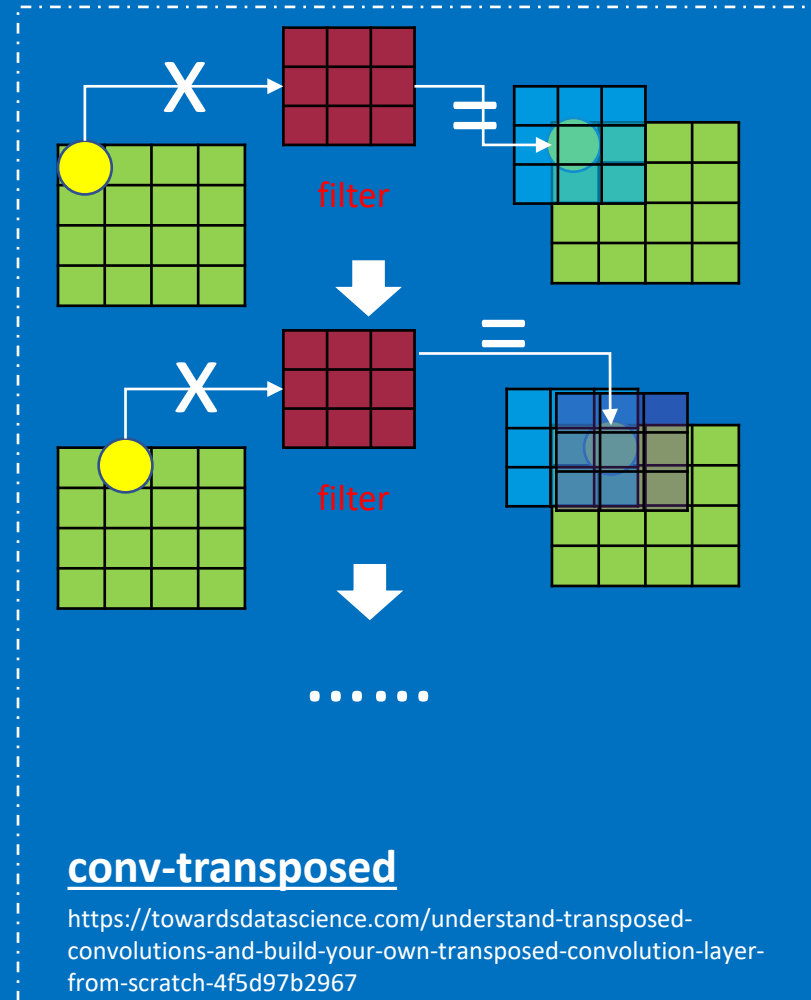
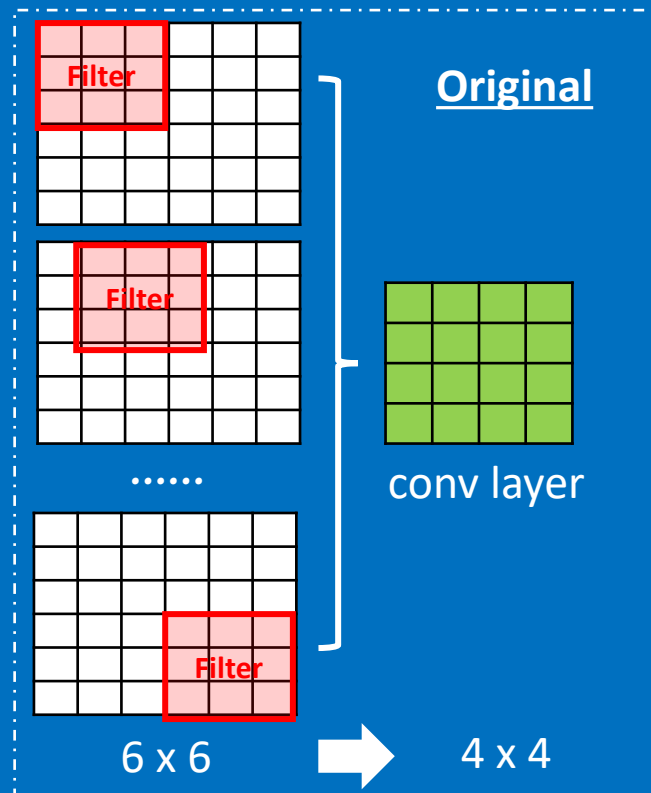
# Best practice for CNN: up-sampling & conv transpose

<https://towardsdatascience.com/understand-transposed-convolutions-and-build-your-own-transposed-convolution-layer-from-scratch-4f5d97b2967>

## What is conv-transposed

In order to resume the shape from

- **convolution**: we can use padding, or **conv-transposed**
- **max-pooling**: we can use up-sampling or conv-transpose



# Best practice for CNN: up-sampling & conv transpose

<https://towardsdatascience.com/understand-transposed-convolutions-and-build-your-own-transposed-convolution-layer-from-scratch-4f5d97b2967>

## What is up-sampling

In order to resume the shape from

- convolution: we can use padding, or conv-tranposed
- **max-pooling**: we can use **up-sampling**

<https://www.machinecurve.com/index.php/2019/12/11/upsampling2d-how-to-use-upsampling-with-keras/#what-is-upsampling>

23	17
24	29

output after max-pooling

23	21.5	18.5	17
23.25	22.44	20.81	20
23.75	24.31	25.43	26
24	25.25	27.75	29

up-sampling

interpolation='bilinear'

23	23	17	17
23	23	17	17
24	24	29	29
24	24	29	29

up-sampling

interpolation='nearest'

# Best practice for CNN: Why Dropout

<https://stackoverflow.com/questions/59717290/does-maxpooling-reduce-overfitting>